



RPNG

Background on Deep Learning

Nate Merrill

University of Delaware

Popular Deep Learning Tasks

- Classification
- Classification + Bounding Box Detection
- Multi-object detection
- Pixel-wise semantic segmentation

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**




CAT, DOG, DUCK

Single object

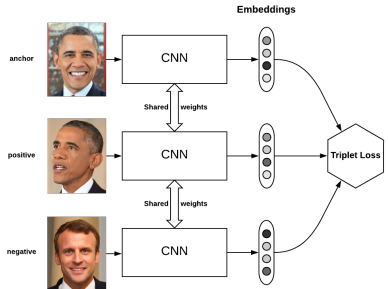
Multiple objects

More Popular Deep Learning Tasks

- Embeddings (dimension reduction)
 - Autoencoder
 - Triplet Embedding



Autoencoder



Triplet embedding

Fully Connected Layer


- Vector input and output
- Learned weights associated with each connection
- Can be written as a linear operation

$$\mathbf{x} = [x_1 \ x_2 \ \dots \ x_n \ 1]^T$$

$$\mathbf{y} = [y_1 \ y_2 \ \dots \ y_m]^T$$

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & \dots & w_{1n} & b_1 \\ w_{21} & w_{22} & \dots & w_{2n} & b_2 \\ \vdots & \ddots & \ddots & \vdots & \vdots \\ w_{m1} & w_{m2} & \dots & w_{mn} & b_m \end{bmatrix}$$




$$\mathbf{y} = \mathbf{W}\mathbf{x}$$



Nonlinear Activation

- Fully connected layer cannot model nonlinear functions
- Nonlinear activations are used to provide nonlinearity
- Key idea: they resemble a "neuron" firing

Examples:



Popular activation functions

Convolution Layer

- Matrix/Tensor input and output
- Useful for image input

X ₁₁	X ₁₂	X ₁₃	X ₁₄	X ₁₅	X ₁₆	X ₁₇
X ₂₁	X ₂₂	X ₂₃	X ₂₄	X ₂₅	X ₂₆	X ₂₇
X ₃₁	X ₃₂	X ₃₃	X ₃₄	X ₃₅	X ₃₆	X ₃₇
X ₄₁	X ₄₂	X ₄₃	X ₄₄	X ₄₅	X ₄₆	X ₄₇
X ₅₁	X ₅₂	X ₅₃	X ₅₄	X ₅₅	X ₅₆	X ₅₇
X ₆₁	X ₆₂	X ₆₃	X ₆₄	X ₆₅	X ₆₆	X ₆₇
X ₇₁	X ₇₂	X ₇₃	X ₇₄	X ₇₅	X ₇₆	X ₇₇

*

K ₁₁	K ₁₂	K ₁₃
K ₂₁	K ₂₂	K ₂₃
K ₃₁	K ₃₂	K ₃₃


=

Y ₁₁	Y ₁₂	Y ₁₃	Y ₁₄	Y ₁₅	Y ₁₆
Y ₂₁	Y ₂₂	Y ₂₃	Y ₂₄	Y ₂₅	Y ₂₆
Y ₃₁	Y ₃₂	Y ₃₃	Y ₃₄	Y ₃₅	Y ₃₆
Y ₄₁	Y ₄₂	Y ₄₃	Y ₄₄	Y ₄₅	Y ₄₆
Y ₅₁	Y ₅₂	Y ₅₃	Y ₅₄	Y ₅₅	Y ₅₆
Y ₆₁	Y ₆₂	Y ₆₃	Y ₆₄	Y ₆₅	Y ₆₆

$$\begin{aligned}
 Y_{ij} = & K_{33}X_{i,j} + K_{32}X_{i,j+1} + K_{31}X_{i,j+2} + K_{23}X_{i+1,j} \\
 & + K_{22}X_{i+1,j+1} + K_{21}X_{i+1,j+2} + K_{13}X_{i+2,j} \\
 & + K_{12}X_{i+2,j+1} + K_{11}X_{i+2,j+2} + b
 \end{aligned}$$

Max Pooling Layer

- Useful for translational invariance
- Similar operation to convolution



Softmax Layer

- Bounds output to $[0,1]$
- Sum of output is 1
- Useful for learning probability mass functions

Suppose $\mathbf{x} \in \mathbb{R}^n$, then

$$\text{Softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

Loss Function

- A scalar error metric for the network
- Necessary for training. Otherwise there is no optimization criteria

Popular examples:

Cross Entropy

Suppose $\mathbf{x} \in \{0, 1\}^n$ is the ground truth class labels and $\hat{\mathbf{x}} \in [0, 1]^n$ s.t. $\sum_{i=1}^n \hat{x}_i = 1$ is the output of a *classification* network, then


$$E(\mathbf{x}, \hat{\mathbf{x}}) = - \sum_{i=1}^n x_i \ln(\hat{x}_i)$$

MSE Suppose $\mathbf{x} \in \mathbb{R}^n$ is the ground truth and $\hat{\mathbf{x}} \in \mathbb{R}^n$ is the output of a *regression* network, then

$$E(\mathbf{x}, \hat{\mathbf{x}}) = \frac{1}{n} \|\hat{\mathbf{x}} - \mathbf{x}\|_2^2$$


Optimization

- Suppose we have a feed-forward neural network with arbitrary tensor input \mathbf{X} , output (loss/error) E , and layers L_1, L_2, \dots, L_k (including loss layer).
- How can we optimize such a complicated function?
- There are many local minima to get caught in



Optimization

- Suppose we have a feed-forward neural network with arbitrary tensor input \mathbf{X} , output (loss/error) E , and layers L_1, L_2, \dots, L_k (including loss layer).
- How can we optimize such a complicated function?
- There are many local minima to get caught in
 - **Idea:** Look at each layer separately
 - Observe that $E = L_k(L_{k-1}(\dots(L_1(\mathbf{X}))))$.
 - We can estimate the gradient of the entire network with the chain rule!



Optimization: Stochastic Gradient Descent

- Create copies of your network
 - Train on multiple data samples at once (i.e., a *batch*)
 - Average out the loss values across examples
 - Make sure the examples are selected *randomly*
- At each iteration, estimate the gradient of each layer w.r.t. the weights
- Update the weights with the *average* gradient over the batch with a scalar multiplier called the *learning rate* (typically $\ll 1$)

Suppose we have the weights of layer L_i , at iteration j , $\mathbf{w}_i^{(j)}$, with layer input \mathbf{X} , learning rate η , and n examples per batch. Then the update step is:

$$\mathbf{w}_i^{(j+1)} = \mathbf{w}_i^{(j)} - \frac{\eta}{n} \sum_{k=1}^n \frac{\partial L_i(\mathbf{w}_i^{(j)}, \mathbf{x}_k)}{\partial \mathbf{w}_i^{(j)}}$$

Implementation

Popular Libraries:

- **TensorFlow**
- Keras
- Caffe (Caffe2 as well)
- PyTorch
- MXNet
- Darknet



The slides and code for the tutorial can be found at:

<https://github.com/nmerrill67/DeepLearningTutorial>

We will step through the script together.