

Implementing an MNIST Variational Autoencoder in TensorFlow

Nate Merrill

October 17, 2019

The MNIST hand-written digits dataset is by far the most famous computer vision dataset available. In this tutorial, we will step through how to create a simple variational autoencoder for this data using the TensorFlow library in python.

The first step is to import the necessary modules. For this script we will need tensorflow, numpy, and matplotlib.pyplot.

```
import tensorflow as tf
import numpy as np
from matplotlib import pyplot as plt
mnist = tf.keras.datasets.mnist
```

Next we need to define the class for our model. We are encoding the input images into a latent space of dimension 2 so that we can visualize the results. In this script, I opted to use three four fully-connected layers to decrease the input dimension to 4, which is then sliced in half to make our two latent variables, μ and σ . The latent variables are trained to parameterize a Gaussian distribution $\mathcal{N}(\mu, \text{diag}(\exp(\sigma)))$. In this implementation, σ is the natural logarithm of the variance for our distribution. The latent variables are optimized to construct a standard normal distribution via the Kullback-Leibler Divergence:

$$\mathcal{D}[\mathcal{N}(\mu, \text{diag}(\exp(\sigma))) || \mathcal{N}(\mathbf{0}, \mathbf{I})] = \frac{1}{2}((\sum_i \exp(\sigma_i) - \sigma_i + \mu^\top \mu - \dim(\mu)) \quad (1)$$

With the latent variables, we can now construct the input to the decoder. ϵ is sampled from a standard normal distribution, and the latent variable $\mathbf{z} = \mu + \text{diag}(\exp(\sigma))^{\frac{1}{2}} \epsilon$ is defined. \mathbf{z} is equivalent to sampling our constructed distribution, but, since sampling is nondifferentiable, we have to resort to what is called the “resampling technique” which uses a sampled standard normal random variable to perform an operation that is equivalent to sampling our distribution, but differentiable.

\mathbf{z} is then fed into the decoder, which uses fully connected layers to upscale up to the original image resolution. The output of the decoder is then sent to the reconstruction loss:

$$- \sum_{h,w,c} x_{h,w,c} \log(r_{h,w,c}) + (1 - x_{h,w,c}) \log(1 - r_{h,w,c}) \quad (2)$$

```

class VAE:
    def __init__(self, epochs=5, batch_sz=16):
        """Constructor"""

        # Input data Tensor
        self.x = tf.placeholder(tf.float32, shape=[None, 28*2])
        fc = self.x
        for n in [256, 128, 64]:
            fc = tf.layers.dense(fc, units=n,
                                activation=tf.nn.elu)
        fc = tf.layers.dense(fc, units=4,
                              activation=None)

        mu = fc[:, :2]

        # estimate log sigma^2 for numerical stability
        log_sig_sq = fc[:, 2:]

        # z = mu + sigma * epsilon
        # epsilon is a sample from a N(0, 1) distribution
        eps = tf.random_normal(tf.shape(mu), 0.0, 1.0, dtype=tf.float32)

        # Random normal variable for decoder :D
        self.z = tf.add(mu, tf.sqrt(tf.exp(log_sig_sq)) * eps, name='z')
        fc = self.z
        for n in [256, 128, 64]:
            fc = tf.layers.dense(fc, units=n,
                                activation=tf.nn.elu)
        fc = tf.layers.dense(fc, units=4,
                              activation=None)

        # sigmoid bounds answer to [0,1] for images
        rec = tf.layers.dense(fc, units=28*2,
                              activation=tf.sigmoid)
        self.rec = rec
        # reconstruction loss
        self.rec_loss = tf.reduce_mean(
            -tf.reduce_sum(self.x * tf.log(tf.clip_by_value(rec, 1e-10, 1.0))
                          + (1.0 - self.x) * tf.log(tf.clip_by_value(1.0 - rec, 1e-10, 1.0))),
            axis=1))

        # stdev is the diagonal of the covariance matrix
        # KLD( (mu, sigma) , N(0,1) ) = .5 (tr(sigma2) + mu^T mu - k - log det(sigma2))

        self.kld = tf.reduce_mean(
            -0.5 * (tf.reduce_sum(1.0 + log_sig_sq - tf.square(mu)
                                - tf.exp(log_sig_sq), axis=1)))

```

```

self.loss = self.kld + self.rec_loss
self.opt = tf.train.AdamOptimizer().minimize(self.loss)

def step(self, batch, sess):
    kld, rec_loss, _ = sess.run([self.kld, self.rec_loss, self.opt],
                                feed_dict={self.x: batch})
    return kld, rec_loss

```

After we train the network, we wish to visualize the results. This function will create a grid over $[-2, 2]$ in each direction and evaluate the decoder at each 2D location. We will see later that the autoencoder distributes each type of digit well throughout the latent space.

```

def display_manifold(model, sess):
    # display a 30x30 2D manifold of digits
    n = 30
    figure = np.zeros((28 * n, 28 * n))
    # linearly spaced coordinates corresponding to the 2D plot
    # of digit classes in the latent space
    grid_x = np.linspace(-2, 2, n)
    grid_y = np.linspace(-2, 2, n)[::-1]

    for i, yi in enumerate(grid_y):
        for j, xi in enumerate(grid_x):
            z_sample = np.array([[xi, yi]])
            x_decoded = sess.run(model.rec,
                                feed_dict={model.z: z_sample})
            digit = x_decoded[0].reshape(28, 28)
            figure[i * 28: (i + 1) * 28,
                  j * 28: (j + 1) * 28] = digit

    plt.figure(figsize=(10, 10))
    start_range = 28 // 2
    end_range = (n - 1) * 28 + start_range + 1
    pixel_range = np.arange(start_range, end_range, 28)
    sample_range_x = np.round(grid_x, 1)
    sample_range_y = np.round(grid_y, 1)
    plt.xticks(pixel_range, sample_range_x)
    plt.yticks(pixel_range, sample_range_y)
    plt.xlabel("z[0]")
    plt.ylabel("z[1]")
    plt.imshow(figure, cmap='Greys_r')
    plt.savefig("latent.png")
    plt.show()

```

Finally, we need to load the MNIST data and train the network with random batches. It is easy to use `numpy.random.randint` to get random indices for the data tensor, and just index them directly. Here I am using a batch size of 256, and training for 5000 iterations. This should be able to run on a CPU in just a few minutes. At the end of training, we can run our `display_manifold` function to visualize the results.

```

if __name__ == '__main__':

    # ignore testing data and labels
    (x_train, _), _ = mnist.load_data()
    x_train = (x_train / 255.0).reshape(-1, 28**2)

    model = VAE()

    batch = 256
    iterations = 5000

    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        for i in range(iterations):
            batch_idx = np.random.randint(len(x_train), size=batch)
            x_batch = x_train[batch_idx, :]


            kld, rec_loss = model.step(x_batch, sess)

            if i % 100 == 0:
                print("Iteration%d:KLD=%f,RecLoss=%f" % (i, kld, rec_loss))

    display_manifold(model, sess)

```

Here you can see what my manifold looks like:



The network is able to encode information about all of the hand-written digits with just two degrees of freedom.